

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

df=pd.read_csv("HousingData.csv")

df

```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222
..
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	NaN	36.2
..
501	21.0	391.99	NaN	22.4
502	21.0	396.90	9.08	20.6
503	21.0	396.90	5.64	23.9
504	21.0	393.45	6.48	22.0
505	21.0	396.90	7.88	11.9

```
[506 rows x 14 columns]
```

```
df.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
PTRATIO \										
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296
15.3										
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242
17.8										
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242
17.8										
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222
18.7										
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222
18.7										

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2

```
df.tail()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
PTRATIO \										
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273
21.0										
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273
21.0										
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273
21.0										
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273
21.0										
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273
21.0										

	B	LSTAT	MEDV
501	391.99	NaN	22.4
502	396.90	9.08	20.6
503	396.90	5.64	23.9
504	393.45	6.48	22.0
505	396.90	7.88	11.9

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
```

```
Data columns (total 14 columns):
 #   Column    Non-Null Count  Dtype  
 --- 
 0   CRIM      486 non-null    float64
 1   ZN        486 non-null    float64
 2   INDUS     486 non-null    float64
 3   CHAS      486 non-null    float64
 4   NOX       506 non-null    float64
 5   RM         506 non-null    float64
 6   AGE        486 non-null    float64
 7   DIS        506 non-null    float64
 8   RAD        506 non-null    int64  
 9   TAX        506 non-null    int64  
 10  PTRATIO   506 non-null    float64
 11  B          506 non-null    float64
 12  LSTAT     486 non-null    float64
 13  MEDV      506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

```
df.describe()
```

	CRIM	ZN	INDUS	CHAS	NOX
RM \ count	486.000000	486.000000	486.000000	486.000000	506.000000
mean	3.611874	11.211934	11.083992	0.069959	0.554695
std	8.720192	23.388876	6.835896	0.255340	0.115878
min	0.006320	0.000000	0.460000	0.000000	0.385000
25%	0.081900	0.000000	5.190000	0.000000	0.449000
50%	0.253715	0.000000	9.690000	0.000000	0.538000
75%	3.560263	12.500000	18.100000	0.000000	0.624000
max	88.976200	100.000000	27.740000	1.000000	0.871000

	AGE	DIS	RAD	TAX	PTRATIO
B \ count	486.000000	506.000000	506.000000	506.000000	506.000000
mean	68.518519	3.795043	9.549407	408.237154	18.455534
std	27.999513	2.105710	8.707259	168.537116	2.164946
min	2.900000	1.129600	1.000000	187.000000	12.600000

```
0.320000
25%    45.175000    2.100175    4.000000    279.000000    17.400000
375.377500
50%    76.800000    3.207450    5.000000    330.000000    19.050000
391.440000
75%    93.975000    5.188425    24.000000    666.000000    20.200000
396.225000
max    100.000000    12.126500    24.000000    711.000000    22.000000
396.900000
```

```
          LSTAT      MEDV
count  486.000000  506.000000
mean   12.715432  22.532806
std    7.155871  9.197104
min    1.730000  5.000000
25%    7.125000  17.025000
50%    11.430000  21.200000
75%    16.955000  25.000000
max    37.970000  50.000000
```

```
df.shape
```

```
(506, 14)
```

```
df.isnull().sum()
```

```
CRIM      20
ZN        20
INDUS     20
CHAS      20
NOX       0
RM        0
AGE       20
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     20
MEDV     0
dtype: int64
```

```
df.columns
```

```
Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS',
'RAD', 'TAX',
'PTRATIO', 'B', 'LSTAT', 'MEDV'],
      dtype='object')
```

```
df.fillna({"CRIM": df["CRIM"].mean(),
           "ZN": df["ZN"].mean(),
```

```

    "INDUS": df["INDUS"].mean(),
    "CHAS": df["CHAS"].mean(),
    "AGE": df["AGE"].mean(),
    "LSTAT": df["LSTAT"].mean()}}, inplace=True)

df.isnull().sum()

CRIM      0
ZN         0
INDUS     0
CHAS      0
NOX       0
RM         0
AGE        0
DIS        0
RAD        0
TAX        0
PTRATIO   0
B          0
LSTAT     0
MEDV      0
dtype: int64

df.corr()

      CRIM      ZN      INDUS      CHAS      NOX      RM
AGE  \
CRIM      1.000000 -0.182930  0.391161 -0.052223  0.410377 -0.215434
0.344934
ZN        -0.182930  1.000000 -0.513336 -0.036147 -0.502287  0.316550 -
0.541274
INDUS     0.391161 -0.513336  1.000000  0.058035  0.740965 -0.381457
0.614592
CHAS      -0.052223 -0.036147  0.058035  1.000000  0.073286  0.102284
0.075206
NOX       0.410377 -0.502287  0.740965  0.073286  1.000000 -0.302188
0.711461
RM        -0.215434  0.316550 -0.381457  0.102284 -0.302188  1.000000 -
0.241351
AGE      0.344934 -0.541274  0.614592  0.075206  0.711461 -0.241351
1.000000
DIS      -0.366523  0.638388 -0.699639 -0.091680 -0.769230  0.205246 -
0.724353
RAD      0.608886 -0.306316  0.593176  0.001425  0.611441 -0.209847
0.449989
TAX      0.566528 -0.308334  0.716062 -0.031483  0.668023 -0.292048
0.500589
PTRATIO  0.273384 -0.403085  0.384806 -0.109310  0.188933 -0.355501
0.262723
B        -0.370163  0.167431 -0.354597  0.050055 -0.380051  0.128069 -

```

```

0.265282
LSTAT    0.434044 -0.407549  0.567354 -0.046166  0.572379 -0.602962
0.574893
MEDV     -0.379695  0.365943 -0.478657  0.179882 -0.427321  0.695360 -
0.380223

          DIS      RAD      TAX      PTRATIO       B      LSTAT
MEDV
CRIM    -0.366523  0.608886  0.566528  0.273384 -0.370163  0.434044 -
0.379695
ZN      0.638388 -0.306316 -0.308334 -0.403085  0.167431 -0.407549
0.365943
INDUS   -0.699639  0.593176  0.716062  0.384806 -0.354597  0.567354 -
0.478657
CHAS    -0.091680  0.001425 -0.031483 -0.109310  0.050055 -0.046166
0.179882
NOX     -0.769230  0.611441  0.668023  0.188933 -0.380051  0.572379 -
0.427321
RM      0.205246 -0.209847 -0.292048 -0.355501  0.128069 -0.602962
0.695360
AGE     -0.724353  0.449989  0.500589  0.262723 -0.265282  0.574893 -
0.380223
DIS     1.000000 -0.494588 -0.534432 -0.232471  0.291512 -0.483429
0.249929
RAD     -0.494588  1.000000  0.910228  0.464741 -0.444413  0.468440 -
0.381626
TAX     -0.534432  0.910228  1.000000  0.460853 -0.441808  0.524545 -
0.468536
PTRATIO -0.232471  0.464741  0.460853  1.000000 -0.177383  0.373343 -
0.507787
B       0.291512 -0.444413 -0.441808 -0.177383  1.000000 -0.368886
0.333461
LSTAT   -0.483429  0.468440  0.524545  0.373343 -0.368886  1.000000 -
0.721975
MEDV    0.249929 -0.381626 -0.468536 -0.507787  0.333461 -0.721975
1.000000

```

Selecting features (columns 0 to 12) and target (column 13)

```

X = df.iloc[:, 0:13]
y = df.iloc[:, 13]

```

```

x_train, x_test, y_train, y_test = train_test_split(X,
y,test_size=0.2,random_state=42)

```

```

print(x_train)

```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
DIS	RAD \						
477	15.02340	0.000000	18.10	0.0	0.6140	5.304	97.300000
2.1007	24						

15	0.62739	0.000000	8.14	0.0	0.5380	5.834	56.500000
4.4986	4						
332	0.03466	11.211934	6.06	0.0	0.4379	6.031	23.300000
6.6407	1						
423	7.05042	0.000000	18.10	0.0	0.6140	6.103	68.518519
2.0218	24						
19	0.72580	0.000000	8.14	0.0	0.5380	5.727	69.500000
3.7965	4						
..
106	0.17120	0.000000	8.56	0.0	0.5200	5.836	91.900000
2.2110	5						
270	0.29916	20.000000	6.96	0.0	0.4640	5.856	42.100000
4.4290	3						
348	0.01501	80.000000	2.01	0.0	0.4350	6.635	29.700000
8.3440	4						
435	11.16040	0.000000	18.10	0.0	0.7400	6.629	94.600000
2.1247	24						
102	0.22876	0.000000	8.56	0.0	0.5200	6.405	85.400000
2.7147	5						
	TAX	PTRATIO	B	LSTAT			
477	666	20.2	349.48	24.91			
15	307	21.0	395.62	8.47			
332	304	16.9	362.25	7.83			
423	666	20.2	2.52	23.29			
19	307	21.0	390.95	11.28			
..			
106	384	20.9	395.67	18.66			
270	223	18.6	388.65	13.00			
348	280	17.0	390.94	5.99			
435	666	20.2	109.85	23.27			
102	384	20.9	70.80	10.63			

[404 rows x 13 columns]

x train.shape

(404, 13)

```
print(x test)
```

72	0.09164	0.0	10.810000	0.0	0.413	6.065	7.800000	5.2873
4								
452	5.09017	0.0	18.100000	0.0	0.713	6.297	91.800000	2.3682
24								
..
412	18.81100	0.0	18.100000	0.0	0.597	4.628	100.000000	1.5539
24								
436	14.42080	0.0	18.100000	0.0	0.740	6.461	93.300000	2.0026
24								
411	14.05070	0.0	18.100000	0.0	0.597	6.657	100.000000	1.5275
24								
86	0.05188	0.0	4.490000	0.0	0.449	6.015	45.100000	4.4272
3								
75	0.09512	0.0	12.830000	0.0	0.437	6.286	45.000000	4.5026
5								

	TAX	PTRATIO	B	LSTAT
173	296	16.6	395.50	9.04
274	254	17.6	396.90	3.53
491	711	20.1	390.11	18.07
72	305	19.2	390.91	5.52
452	666	20.2	385.09	17.27
..
412	666	20.2	28.79	34.37
436	666	20.2	27.49	18.05
411	666	20.2	35.05	21.22
86	247	18.5	395.99	12.86
75	398	18.7	383.23	8.94

[102 rows x 13 columns]

x_test.shape

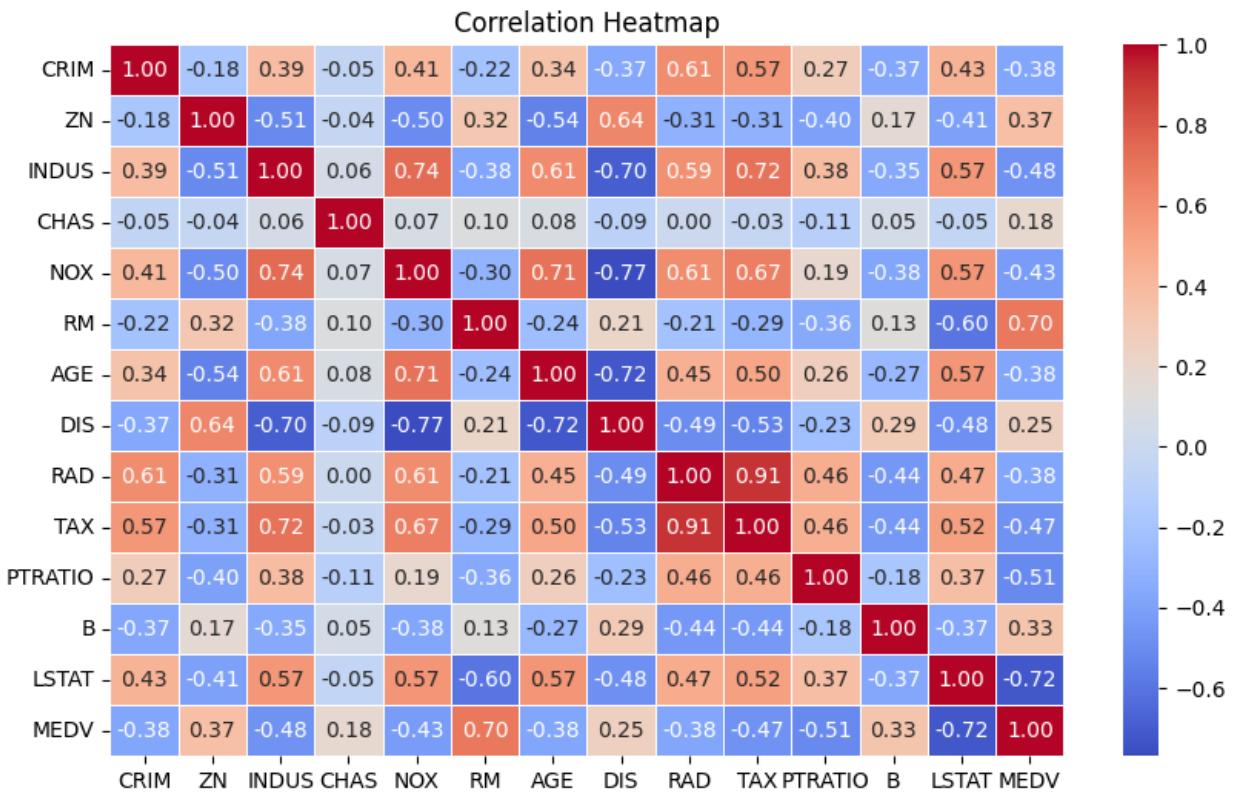
(102, 13)

print(y_train)

477	12.0
15	19.9
332	19.4
423	13.4
19	18.2
..	..
106	19.5
270	21.1
348	24.5
435	13.4
102	18.6

Name: MEDV, Length: 404, dtype: float64

```
y_train.shape  
(404,)  
  
print(y_test)  
  
173    23.6  
274    32.4  
491    13.6  
72     22.8  
452    16.1  
...  
412    17.9  
436     9.6  
411    17.2  
86     22.5  
75     21.4  
Name: MEDV, Length: 102, dtype: float64  
  
y_test.shape  
(102,)  
  
corr_matrix = df.corr()  
  
# Set up the heatmap  
plt.figure(figsize=(10, 6)) # Adjust figure size  
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", fmt=".2f",  
lineweights=0.5)  
  
# Show the plot  
plt.title("Correlation Heatmap")  
plt.show()
```



```

linear_regression=LinearRegression()
linear_regression.fit(x_train,y_train)
LinearRegression()
y_pred=linear_regression.predict(x_test)

mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print(f"Mean Absolute Error: {mae}")
print(f"Mean Squared Error: {mse}")
print(f"Root Mean Squared Error: {rmse}")
print(f"R-squared Score: {r2}")

Mean Absolute Error: 3.149923357345782
Mean Squared Error: 25.017672023842703
Root Mean Squared Error: 5.001766890194174
R-squared Score: 0.658852019550814

# Visualization: Actual vs. Predicted Prices
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred)

```

```
plt.xlabel("Actual ")
plt.ylabel("Predicted ")
plt.title("Actual vs. Predicted Prices")
plt.show()
```

