

**reducer.py**

```

import sys

# Initialize variables
temperature_sum = 0
temperature_count = 0
dewpoint_sum = 0
dewpoint_count = 0
windspeed_sum = 0
windspeed_count = 0

# Read input from standard input
for line in sys.stdin:
    # Remove leading and trailing whitespace
    line = line.strip()

    # Split the line into key and value
    key, value = line.split('\t', 1)

    # Update sums and counts based on key
    if key == 'temperature':
        temperature_sum += float(value)
        temperature_count += 1
    elif key == 'dewpoint':
        dewpoint_sum += float(value)
        dewpoint_count += 1
    elif key == 'windspeed':
        windspeed_sum += float(value)
        windspeed_count += 1

# Calculate averages
average_temperature = temperature_sum / temperature_count if temperature_count > 0 else 0
average_dewpoint = dewpoint_sum / dewpoint_count if dewpoint_count > 0 else 0
average_windspeed = windspeed_sum / windspeed_count if windspeed_count > 0 else 0

# Output the averages
print('Average Temperature:', average_temperature)
print('Average Dew Point:', average_dewpoint)
print('Average Wind Speed:', average_windspeed)

```

**mapper.py**

```

import sys

# Initialize variables
temperature_sum = 0
temperature_count = 0

```

```

dewpoint_sum = 0
dewpoint_count = 0
windspeed_sum = 0
windspeed_count = 0

# Read input from standard input
for line in sys.stdin:
    # Remove leading and trailing whitespace
    line = line.strip()

    # Skip empty lines
    if not line:
        continue

    # Split the line into temperature, dew point, and wind speed
    values = line.split(',')

    # Check if all values are present
    if len(values) != 3:
        continue

    try:
        temperature, dewpoint, windspeed = map(float, values)

        # Output key-value pairs
        print('temperature\t{}'.format(temperature))
        print('dewpoint\t{}'.format(dewpoint))
        print('windspeed\t{}'.format(windspeed))
    except ValueError:
        # Handle the case where conversion to float fails
        continue

```

**Terminal:**

```

aiml@aiml-System-Product-Name:~$ cd Desktop/reduce2
aiml@aiml-System-Product-Name:~/Desktop/reduce2$ cat sample_weather.txt | python3 mapper.py
Traceback (most recent call last):
  File "/home/aiml/Desktop/reduce2/mapper.py", line 18, in <module>
    count = int(values_count[1])
IndexError: list index out of range
aiml@aiml-System-Product-Name:~/Desktop/reduce2$ ^C
aiml@aiml-System-Product-Name:~/Desktop/reduce2$ cat sample_weather.txt | python3 mapper.py
Average Temperature: 0
Average Dew Point: 0

```

Average Wind Speed: 0

aiml@aiml-System-Product-Name:~/Desktop/reduce2\$ cat sample\_weather.txt | python3 mapper.py

Traceback (most recent call last):

File "/home/aiml/Desktop/reduce2/mapper.py", line 17, in <module>

    temperature, dewpoint, windspeed = map(float, line.split(','))

ValueError: could not convert string to float: "

aiml@aiml-System-Product-Name:~/Desktop/reduce2\$ cat sample\_weather.txt | python3 mapper.py

temperature\_sum: 129.1

temperature\_count: 6

dewpoint\_sum: 81.3999999999999

dewpoint\_count: 6

windspeed\_sum: 35.7

windspeed\_count: 6

aiml@aiml-System-Product-Name:~/Desktop/reduce2\$ cat sample\_weather.txt | python3 mapper.py

Traceback (most recent call last):

File "/home/aiml/Desktop/reduce2/mapper.py", line 17, in <module>

    temperature, dewpoint, windspeed = map(float, line.split(','))

ValueError: could not convert string to float: "

aiml@aiml-System-Product-Name:~/Desktop/reduce2\$ cat sample\_weather.txt | python3 mapper.py

Traceback (most recent call last):

File "/home/aiml/Desktop/reduce2/mapper.py", line 17, in <module>

    temperature, dewpoint, windspeed = map(float, line.split(','))

ValueError: could not convert string to float: "

aiml@aiml-System-Product-Name:~/Desktop/reduce2\$ cat sample\_weather.txt | python3 mapper.py

Average Temperature: 21.51666666666666

Average Dew Point: 13.56666666666666

Average Wind Speed: 5.95

aiml@aiml-System-Product-Name:~/Desktop/reduce2\$ cat word\_count\_data.txt | python3

mapper.py | sort -k1,1 | python3 reducer.py

cat: word\_count\_data.txt: No such file or directory

Traceback (most recent call last):

File "/home/aiml/Desktop/reduce2/reducer.py", line 16, in <module>

    key, value = line.split('\t', 1)

ValueError: not enough values to unpack (expected 2, got 1)

aiml@aiml-System-Product-Name:~/Desktop/reduce2\$ cat sample\_weather.txt | python3 mapper.py

temperature 23.4

dewpoint 15.2

windspeed 5.6

temperature 20.1

dewpoint 12.5

```
windspeed    6.3
temperature   18.5
dewpoint     10.8
windspeed    7.2
temperature   25.0
dewpoint     16.9
windspeed    4.8
temperature   22.3
dewpoint     14.7
windspeed    5.1
temperature   19.8
dewpoint     11.3
windspeed    6.7
```

```
aiml@aiml-System-Product-Name:~/Desktop/reduce2$ cat word_count_data.txt | python3
mapper.py | sort -k1,1 | python3 reducer.py
```

```
cat: word_count_data.txt: No such file or directory
```

```
Average Temperature: 0
```

```
Average Dew Point: 0
```

```
Average Wind Speed: 0
```

```
aiml@aiml-System-Product-Name:~/Desktop/reduce2$ cat sample_weather.txt | python3 mapper.py
| sort -k1,1 | python3 reducer.py
```

```
Average Temperature: 21.516666666666666
```

```
Average Dew Point: 13.566666666666668
```

```
Average Wind Speed: 5.949999999999999
```

```
aiml@aiml-System-Product-Name:~/Desktop/reduce2$ cd
```

```
aiml@aiml-System-Product-Name:~$ start-dfs.sh
```

```
Starting namenodes on [localhost]
```

```
localhost: namenode is running as process 28316. Stop it first and ensure /tmp/hadoop-aiml-
namenode.pid file is empty before retry.
```

```
Starting datanodes
```

```
localhost: datanode is running as process 28494. Stop it first and ensure /tmp/hadoop-aiml-
datanode.pid file is empty before retry.
```

```
Starting secondary namenodes [aiml-System-Product-Name]
```

```
aiml-System-Product-Name: secondarynamenode is running as process 28720. Stop it first and
ensure /tmp/hadoop-aiml-secondarynamenode.pid file is empty before retry.
```

```
2024-04-04 11:47:12,356 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
```

```
aiml@aiml-System-Product-Name:~$ start-yarn.sh
```

```
Starting resourcemanager
```

```
resourcemanager is running as process 28950. Stop it first and ensure /tmp/hadoop-aiml-
resourcemanager.pid file is empty before retry.
```

### DSBDA Group B Assignment 3

Starting nodemanagers

localhost: nodemanager is running as process 29291. Stop it first and ensure /tmp/hadoop-aiml-nodemanager.pid file is empty before retry.

```
aiml@aiml-System-Product-Name:~$ hdfs dfs -mkdir /avg_temp_in_python
```

2024-04-04 11:47:53,637 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

```
aiml@aiml-System-Product-Name:~$ hdfs dfs -mkdir /avg_temp_in_python
```

2024-04-04 11:47:58,538 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

```
mkdir: '/avg_temp_in_python': File exists
```

```
aiml@aiml-System-Product-Name:~$ hdfs dfs -copyFromLocal
```

```
/home/aiml/Desktop/reduce2/sample_weather.txt /avg_temp_in_python
```

2024-04-04 11:48:45,619 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

```
aiml@aiml-System-Product-Name:~$ hdfs dfs -ls /
```

2024-04-04 11:49:00,561 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 10 items

drwxr-xr-x - aiml supergroup	0 2024-04-04 11:48 /avg_temp_in_python
-rw-r--r-- 1 aiml supergroup	399 2024-03-15 10:57 /data
-rw-r--r-- 1 aiml supergroup	399 2024-03-15 10:57 /data.txt
drwxr-xr-x - aiml supergroup	0 2024-03-15 10:58 /tmp
drwxr-xr-x - aiml supergroup	0 2024-03-15 10:58 /user
drwxr-xr-x - aiml supergroup	0 2024-04-04 10:06 /weather_forcast_in_python
drwxr-xr-x - aiml supergroup	0 2024-04-03 15:45 /word_count_in_python
drwxr-xr-x - aiml supergroup	0 2024-04-04 10:16 /word_count_in_python2
drwxr-xr-x - aiml supergroup	0 2024-04-03 16:32 /word_count_in_python3
drwxr-xr-x - aiml supergroup	0 2024-04-04 11:25 /word_count_in_python8

```
aiml@aiml-System-Product-Name:~$ hdfs dfs -ls /avg_temp_in_python
```

2024-04-04 11:49:34,286 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 1 items

```
-rw-r--r-- 1 aiml supergroup 85 2024-04-04 11:48 /avg_temp_in_python/sample_weather.txt
```

```
aiml@aiml-System-Product-Name:~$ cd/Desktop/reduce2
```

bash: cd/Desktop/reduce2: No such file or directory

```
aiml@aiml-System-Product-Name:~$ cd/Desktop
```

bash: cd/Desktop: No such file or directory

```
aiml@aiml-System-Product-Name:~$ cd Desktop/reduce2
```

```
aiml@aiml-System-Product-Name:~/Desktop/reduce2$ chmod 777 mapper.py reducer.py
```

```
aiml@aiml-System-Product-Name:~/Desktop/reduce2$ cd
```

### DSBDA Group B Assignment 3

```
aiml@aiml-System-Product-Name:~$ hadoop jar /home/aiml/Desktop/reduce2/hadoop-streaming-3.3.6.jar \
> -input /avg_temp_in_python/sample_weather.txt \
> -output /avg_temp_in_python/output \
> -mapper "python3 /home/aiml/Desktop/reduce2/mapper.py" \
> -reducer "python3 /home/aiml/Desktop/reduce2/reducer.py"
2024-04-04 11:53:18,540 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
packageJobJar: [/tmp/hadoop-unjar1566709327473520848/] []
/tmp/streamjob1191938198619220873.jar tmpDir=null
2024-04-04 11:53:18,872 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /127.0.0.1:8032
2024-04-04 11:53:18,948 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /127.0.0.1:8032
2024-04-04 11:53:19,058 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/aiml/.staging/job_1712204091628_0006
2024-04-04 11:53:19,189 INFO mapred.FileInputFormat: Total input files to process : 1
2024-04-04 11:53:19,211 INFO mapreduce.JobSubmitter: number of splits:2
2024-04-04 11:53:19,268 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1712204091628_0006
2024-04-04 11:53:19,268 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-04-04 11:53:19,364 INFO conf.Configuration: resource-types.xml not found
2024-04-04 11:53:19,364 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-04-04 11:53:19,401 INFO impl.YarnClientImpl: Submitted application application_1712204091628_0006
2024-04-04 11:53:19,422 INFO mapreduce.Job: The url to track the job: http://aiml-System-ProductName:8088/proxy/application_1712204091628_0006/
2024-04-04 11:53:19,423 INFO mapreduce.Job: Running job: job_1712204091628_0006
2024-04-04 11:53:23,468 INFO mapreduce.Job: Job job_1712204091628_0006 running in uber mode : false
2024-04-04 11:53:23,470 INFO mapreduce.Job: map 0% reduce 0%
2024-04-04 11:53:26,520 INFO mapreduce.Job: map 100% reduce 0%
2024-04-04 11:53:30,554 INFO mapreduce.Job: map 100% reduce 100%
2024-04-04 11:53:30,568 INFO mapreduce.Job: Job job_1712204091628_0006 completed successfully
2024-04-04 11:53:30,621 INFO mapreduce.Job: Counters: 54
    File System Counters
        FILE: Number of bytes read=312
        FILE: Number of bytes written=834992
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
```

## DSBDA Group B Assignment 3

HDFS: Number of bytes read=350  
HDFS: Number of bytes written=106  
HDFS: Number of read operations=11  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=2  
HDFS: Number of bytes read erasure-coded=0

### Job Counters

Launched map tasks=2  
Launched reduce tasks=1  
Data-local map tasks=2  
Total time spent by all maps in occupied slots (ms)=2710  
Total time spent by all reduces in occupied slots (ms)=1156  
Total time spent by all map tasks (ms)=2710  
Total time spent by all reduce tasks (ms)=1156  
Total vcore-milliseconds taken by all map tasks=2710  
Total vcore-milliseconds taken by all reduce tasks=1156  
Total megabyte-milliseconds taken by all map tasks=2775040  
Total megabyte-milliseconds taken by all reduce tasks=1183744

### Map-Reduce Framework

Map input records=7  
Map output records=18  
Map output bytes=270  
Map output materialized bytes=318  
Input split bytes=222  
Combine input records=0  
Combine output records=0  
Reduce input groups=3  
Reduce shuffle bytes=318  
Reduce input records=18  
Reduce output records=3  
Spilled Records=36  
Shuffled Maps =2  
Failed Shuffles=0  
Merged Map outputs=2  
GC time elapsed (ms)=89  
CPU time spent (ms)=1010  
Physical memory (bytes) snapshot=870002688  
Virtual memory (bytes) snapshot=7719759872  
Total committed heap usage (bytes)=931135488  
Peak Map Physical memory (bytes)=329068544

### DSBDA Group B Assignment 3

Peak Map Virtual memory (bytes)=2571878400  
Peak Reduce Physical memory (bytes)=216350720  
Peak Reduce Virtual memory (bytes)=2577326080

#### Shuffle Errors

BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0  
WRONG\_MAP=0  
WRONG\_REDUCE=0

#### File Input Format Counters

Bytes Read=128

#### File Output Format Counters

Bytes Written=106

2024-04-04 11:53:30,621 INFO streaming.StreamJob: Output directory:

/avg\_temp\_in\_python/output

aiml@aiml-System-Product-Name:~\$ hdfs dfs -cat /avg\_temp\_in\_python/output/part-00000

2024-04-04 11:54:05,549 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Average Temperature: 21.516666666666666

Average Dew Point: 13.566666666666668

Average Wind Speed: 5.95

aiml@aiml-System-Product-Name:~\$

#### **sample\_weather.txt:**

23.4,15.2,5.6  
20.1,12.5,6.3  
18.5,10.8,7.2  
25.0,16.9,4.8  
22.3,14.7,5.1  
19.8,11.3,6.7